

Convolutional Neural Networks (CNNs)

Theory, Mathematics, and Architecture Design

Prof. D.Sc. BARSEKH-ONJI Aboud

Facultad de Ingeniería
Universidad Anáhuac México

January 21, 2026

Agenda

1 Motivation

2 The Convolution Operation

3 Hyperparameters & Arithmetic

4 CNN Building Blocks

5 MATLAB Implementation

6 Conclusion

Biological Inspiration

CNNs are biologically inspired by the visual cortex organization (Hubel & Wiesel, 1959).

- Local Receptive Fields: Neurons respond only to a small sub region of the visual field.
- Hierarchy:
- **V1 Area:** Detects simple edges and orientations.
- **Higher Areas:** Detect shapes, textures, and objects.

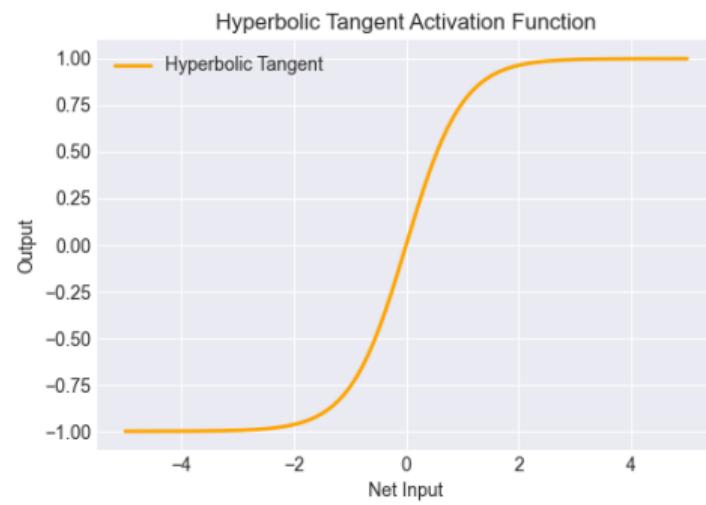


Figure 1: Hierarchical processing in the brain vs. Deep Learning.

Why not Standard MLPs?

Imagine an image of 224×224 pixels (RGB).

The Curse of Dimensionality

If we flatten this image into a vector:

$$\text{Input Vector} = 224 \times 224 \times 3 = 150,528 \text{ inputs}$$

If the first hidden layer has 1,000 neurons (Dense Layer):

$$\text{Weights} = 150,528 \times 1,000 \approx 150 \text{ Million Parameters!}$$

Problem: Huge computational cost and massive overfitting.

Solution: Local connectivity and Parameter Sharing (CNNs).

Agenda

1 Motivation

2 The Convolution Operation

3 Hyperparameters & Arithmetic

4 CNN Building Blocks

5 MATLAB Implementation

6 Conclusion

The Mathematical Definition

In Deep Learning, a "convolution" is technically a **sliding dot product** (cross-correlation).

We pass a filter (kernel) K over an image I :

$$(I * K)(i, j) = \sum_m \sum_n I(i + m, j + n) \cdot K(m, n) \quad (1)$$

- **Element-wise multiplication:** Between the kernel weights and the image patch.
- **Summation:** Aggregating the result into a single pixel in the Feature Map.

Feature Maps & Channels

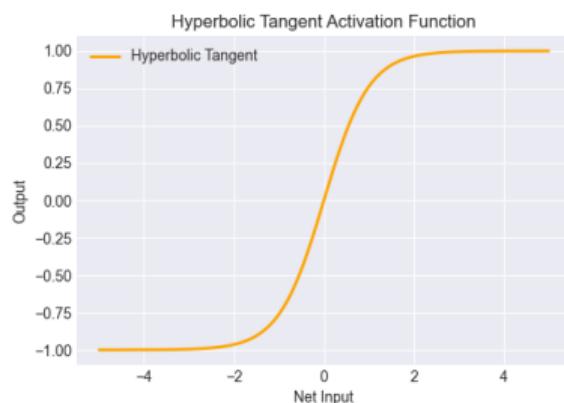
Input Tensor ($H \times W \times C_{in}$)

An image usually has 3 channels (RGB).

Output Tensor ($H' \times W' \times K$)

If we apply K distinct filters (e.g., 32 filters):

- We obtain 32 discrete 2D maps.
- These stack together to form the output depth (Channels).
- **Analogy:** Each channel represents a specific "feature" (one for horizontal lines, one for



[Placeholder: Animation of a kernel sliding over a matrix]

Figure 2: The sliding window

Agenda

1 Motivation

2 The Convolution Operation

3 Hyperparameters & Arithmetic

4 CNN Building Blocks

5 MATLAB Implementation

6 Conclusion

Spatial Dimensions Control

How do we calculate the output size of a layer? This is the fundamental arithmetic equation for CNN design.

$$O = \left\lfloor \frac{W - F + 2P}{S} \right\rfloor + 1 \quad (2)$$

Where:

- W : Input Volume Size (Width/Height)
- F : Filter /Kernel Size (e.g., 3×3)
- S : **Stride** (Step size of the slide)
- P : **Padding** (Zeros added around the border)

Padding Options

1. Valid Padding ($P = 0$):

- No padding.
- The image shrinks after convolution.
- Losing edge information.

2. Same Padding ($P > 0$):

- Padding added to preserve dimensions.
- Output Size = Input Size (if stride=1).
- Allows for very deep networks without vanishing spatial size.

MATLAB syntax: "Padding", "same" vs "Padding", 0

Agenda

1 Motivation

2 The Convolution Operation

3 Hyperparameters & Arithmetic

4 CNN Building Blocks

5 MATLAB Implementation

6 Conclusion

The Non-Linearity (ReLU)

Convolution is a *linear* operation (multiplication + sum). To learn complex patterns, we need non-linearity.

$$f(x) = \max(0, x) \quad (3)$$

Rectified Linear Unit (ReLU):

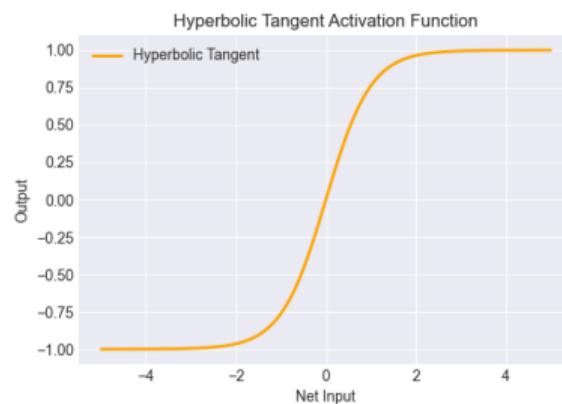
- Introduces sparsity.
- Avoids Vanishing Gradient problem (common in Sigmoid).
- Extremely computationally efficient.

Pooling Layers (Subsampling)

Purpose: Progressively reduce the spatial size (S -dimension) to reduce parameters and computation.

- **Max Pooling:** Takes the maximum value in the window. Extracts the most prominent features.
- **Average Pooling:** Takes the average. Smoothes features.

Key Property: *Translation Invariance*. Small shifts in the image don't change the outcome.



*Placeholder: 2x2 Max Pooling on
a 4x4 matrix*

Flattening & Fully Connected

Transition to Classification

After extracting spatial features with Convolution+Pooling, we get a 3D volume (e.g., $7 \times 7 \times 512$).

- 1 **Flatten:** Unroll the volume into a 1D vector (size 25,088).
- 2 **Fully Connected (FC):** Standard MLP layers to combine features and perform classification logic.
- 3 **Softmax:** Probability distribution for the output classes.

Agenda

1 Motivation

2 The Convolution Operation

3 Hyperparameters & Arithmetic

4 CNN Building Blocks

5 MATLAB Implementation

6 Conclusion

Defining a CNN in MATLAB

Using the Layer Graph approach (standard for modern image classification):

Listing 1: Standard VGG-Style Block

```
1 inputSize = [28 28 1]; % Grayscale image (S-S-C)
2
3 layers = [
4     imageInputLayer(inputSize, "Name", "imageinput")
5
6     % Convolutional Block 1
7     convolution2dLayer(3, 8, "Padding", "same", "Name", "conv_1")
8     batchNormalizationLayer("Name", "batchnorm_1")
9     reluLayer("Name", "relu_1")
10
11    % Downsampling
12    maxPooling2dLayer(2, "Stride", 2, "Name", "maxpool_1")
13
14    % Classification Head
15    fullyConnectedLayer(10, "Name", "fc") % 10 classes
```



Understanding Data Formats: SSCB

When training CNNs, MATLAB (since R2023a) usually expects tensors in the format **SSCB**.

- **S (Spatial)**: Height (H).
- **S (Spatial)**: Width (W).
- **C (Channels)**: Depth (Colors or Feature Maps).
- **B (Batch)**: Number of images processed in parallel.

$$\text{Tensor} = 224 \times 224 \times 3 \times 32$$

Agenda

1 Motivation

2 The Convolution Operation

3 Hyperparameters & Arithmetic

4 CNN Building Blocks

5 MATLAB Implementation

6 Conclusion

Summary

- **Convolution** exploits spatial hierarchies and shared weights, making image processing efficient.
- **ReLU** provides the necessary non-linearity for deep learning.
- **Pooling** reduces dimensionality and grants translation invariance.
- **Architecture:** It is a game of managing Dimensions (S) and Channels (C) usually decreasing spatial size while increasing channel depth deep in the network.